

Gitoci

Git remote helper for storing repositories in OCI


By: Nathan Joslin

Advisor: Dr. Kyle Tarplee




Problem Description

- Need for moving Git repos across air-gapped environments
- Core problems:
 - Efficient data transfer
 - Efficient data storage
 - Minimal maintenance
- Existing tools do not solve all three well
- How can we achieve the main goal by using their strengths?



Problem Description

- Git bundles
 - Good for offline transfers
- ORAS
 - Efficient transfers of OCI artifacts
- How can we transfer Git repositories to/from OCI registries?



Existing Solutions

Zarf

- CD for air gapped systems
- Large tool and “ecosystem”
- Lacks incremental updates
- Lacks git-lfs support



ORAS

- Efficient transfers of OCI artifacts
- Go package
- Command line interface
- CNCF Sandbox project



Git Bundles

- Native to git
- Intended for offline transfers
- “Thin” packs
- Difficult to manage at scale, large and/or many repositories
- Lacks git-lfs support



Proof-of-Concept

ASCE Data Tool:

- Prototype OCI data model using Git bundles
- Existing user base

Pros:

- Supports incremental updates
- Supports updating subsets of all repository references
- All work is done client side
- Supports git-lfs

Cons:

- Room for optimizations
- Subcommand of a larger tool
- Does not support native Git usage

```
$ export IMAGE_NAME=torvalds-linux-$(uuidgen)
$ ace-dt git to-oci git@github.com:torvalds/linux.git ttl.sh/${IMAGE_NAME}:1h
$ ace-dt git from-oci ttl.sh/${IMAGE_NAME} git@github.com:user-example.git
```

Proposed Solution

OCI Spec for Git Repositories + Git Remote Helper

Why OCI?

- An industry standard for image storage and transfer
- Increasing usage for alternative artifacts
- Take advantage of existing transfer tools

Why Git Remote Helper?

- Allow Git to interface with alternative repo formats
- Use git natively, or with existing tools (e.g. VSCode)
- Well-defined requirements



Proposed Solution

- OCI data model using Git packfiles
- Git remote helper to facilitate pushing and fetching to/from OCI
- All work done client side

```
# Clone directly from OCI registry
$ git clone oci://example.registry.com/repository/name:tag

# Add an OCI registry as a remote
$ git remote add <remote-name> oci://example.registry.com/repository/name:tag
```

```
# Fetch from a remote OCI registry
$ git fetch [<repository>] [<refspec>...]

# Push to a remote OCI registry
$ git push [<repository>] [<refspec>...]
```

Goals & Implementation

Goals:

- OCI Spec for Git Repos using packfiles
- Git Remote Helper
- Support for git-lfs
- Minimize overhead compared to native Git
- 50% test coverage

Implementation Details:

- Written in Go
- Key External Dependencies:
 - ORAS
 - Sourcegraph/conc
 - Go-git

Remote Helper Capabilities

Pushing

Capability	Support
connect	Yes
stateless-connect	No
push	Maybe
export	Maybe
no-private-update	No

Fetching

Capability	Support
connect	Yes
stateless-connect	No
fetch	Maybe
import	Maybe
check-connectivity	Maybe
get	Maybe

Misc.

Capability	Support
option	Maybe
refspec	Maybe
bidi-import	Maybe
export-marks	Maybe
import-marks	Maybe
signed-tags	Maybe
object-format	Maybe

Capabilities: Pushing

- Connect (Yes)
 - Utilizes Git's packfile protocol.
 - Bundles are an extension of packfiles, used by proof-of-concept.
 - Necessary for efficient pushes, and OCI storage.
 - Stateless-Connect (No)
 - Described as experimental and intended for internal use.
 - Push (Maybe)
 - Used to update remote history and refs with local objects.
 - Packfile protocol supported by *connect* is better suited for goals.
 - *Connect* is preferred, unless *push* or *export* is requested by remote helper.
 - May improve efficiency in certain situations, experimentation needed.
-

Capabilities: Pushing

- Export (Maybe)
 - Used to push objects over a fast-import stream.
 - Git prefers *push* over *export*.
 - More research needed.
 - No-private-update (No)
 - Adds support for disabling private namespace updates.
 - Not within scope of project.
 - See [Capabilities for Pushing](#)
-

Capabilities: Fetching

- Connect (Yes)
 - Utilizes Git's packfile protocol.
 - Bundles are an extension of packfiles, used by proof-of-concept.
 - Necessary for efficient fetches, and OCI storage.
 - Stateless-Connect (No)
 - Described as experimental and intended for internal use.
 - Fetch (Maybe)
 - Used to update local history and refs with remote objects.
 - Packfile protocol supported by *connect* is better suited for goals.
 - May improve efficiency in certain situations, experimentation needed.
-

Capabilities: Fetching

- Import (Maybe)
 - Used to fetch git objects over fast-import stream from a remote.
 - Maybe be useful to receive packfiles from OCI, and send objects to git.
 - More research necessary.
 - Check-connectivity (Maybe)
 - Used to validate that a packfile, from a clone, is self contained.
 - Not necessary, but potentially useful.
 - Get (Maybe)
 - Used to fetch files from a URI.
 - Not necessary for MVP, but potentially useful.
 - See Fetching Capabilities
-

Capabilities: Miscellaneous

- Option (Maybe)
 - Potentially helpful for controlling verbosity and history depth.
 - Refspec (Maybe)
 - Required for *export*.
 - Optional for *import*.
 - Adds support for constraining references to private namespaces.
 - Bidi-import (Maybe)
 - Modifies *import*, for potential efficiency improvements.
 - “The fast-import commands *cat-blob* and *ls* can be used by remote-helpers to retrieve information about blobs and trees that already exist in fast-import’s memory”.
 - Export-marks (Maybe)
 - Modifiest *export*.
 - Git provides output of internal marks table.
 - Potential efficiency improvements.
-

Capabilities: Miscellaneous

- Import-marks (Maybe)
 - More information needed.
 - Intuitively one would think it's the reverse of *export-marks*, but not according to Git's docs.
 - Signed-tags (Maybe)
 - Modifies *export* capability regarding how signed tags are passed to *git-fast-export*.
 - Looking into how Git handles tags in *fast-import* and *fast-export* is worthwhile.
 - Object-format (Maybe)
 - Indicates the helper can use explicit hash algorithm extensions when interacting with remote.
 - Potentially useful for supporting alternative hash algorithms, as done by OCI.
 - See Miscellaneous Capabilities
-

Questions?



Sources

- Git logos by [Jason Long](#) - [CC BY 3.0](#)
- Open Container Initiative logos obtained from [GitHub](#) - [Apache 2.0](#)
- [Zarf logo](#) - [Apache 2.0](#)
- [Git remote helpers](#)
- [ORAS](#)
- [ASCE Data Tool](#)